


CNIT 127: Exploit Development

Ch 2: Stack Overflows in Linux



Updated 8-28-19

Topics

- Buffers in C
- Information Disclosure
- gdb: Gnu Debugger
- Segmentation Fault
- The Stack
- Functions and the Stack
- Stack Buffer Overflow

Stack-Based Buffer Overflows

- Most popular and best understood exploitation method
- Aleph One's "Smashing the Stack for Fun and Profit" (1996)
 - Link Ch 2a
- Buffer
 - A limited, contiguously allocated set of memory
 - In C, usually an *array*

Exploit A: Information Disclosure

C and C++ Lack Bounds-Checking

- It is the programmer's responsibility to ensure that array indices remain in the valid range

```
#include <stdio.h>

int main()
{
    int array[5] = {1, 2, 3, 4, 5};
    printf("%d\n", array[5]);
}
```

Reading Past End of Array

```
GNU nano 2.9.2 ch2a.c

#include <stdio.h>

int main()
{
    int array[5] = {1, 2, 3, 4, 5};
    printf("%d\n", array[5]);
}
```

```
cnitfiftythree@deb:~/127/ch2$ gcc -m32 -g -o ch2a ch2a.c
cnitfiftythree@deb:~/127/ch2$ ./ch2a
-10544
```

- We can read data that we shouldn't be seeing
- *Information disclosure vulnerability*

Using gdb (GNU Debugger)

```
[root@kali:~/127/ch2# gdb -q ch2a
Reading symbols from ch2a...done.
(gdb) list
1      #include <stdio.h>
2
3      int main()
4      {
5          int array[5] = {1, 2, 3, 4, 5};
6          printf("%d\n", array[5]);
7      }
8
(gdb)
```

- Source code debugging
- Because we compiled with `gcc -g`

Using gdb (GNU Debugger)

```
(gdb) break 6
Breakpoint 1 at 0x5df: file ch2a.c, line 6.
(gdb) run
Starting program: /home/cnitfiftythree/127/ch2/ch2a

Breakpoint 1, main () at ch2a.c:6
6          printf("%d\n", array[5]);
(gdb) x/8x &array
0xffffd63c:    0x00000001    0x00000002    0x00000003    0x00000004
0xffffd64c:    0x00000005    0xffffd670    0x00000000    0x00000000
(gdb) █
```

- gdb commands

list	<i>show source code</i>
run	<i>execute program</i>
break	<i>insert breakpoint</i>
x	<i>examine memory</i>

Exploit B: Denial of Service

Writing Past End of Array

```
GNU nano 2.9.2 ch2b.c

#include <stdio.h>

int main()
{
    int array[5];
    array[1000] = 1;
}
```

```
cnitfiftythree@deb:~/127/ch2$ gcc -m32 -g -o ch2b ch2b.c
cnitfiftythree@deb:~/127/ch2$ ./ch2b
Segmentation fault
cnitfiftythree@deb:~/127/ch2$
```

- Program has crashed
- *Denial of service*

Print Out More Information

```
GNU nano 2.9.2 ch2c.c

#include <stdio.h>

int main()
{
    int array[5], i;

    for (i=0; i<=1000; i+=100)
    {
        printf("%x\n", &array[i]);
        array[i] = 1;
    }
}
```

- printf uses a **format string**
- **%x** means print in hexadecimal

Segmentation Fault

```
cnitfiftythree@deb:~/127/ch2$ gcc -m32 -g -o ch2c ch2c.c
cnitfiftythree@deb:~/127/ch2$ ./ch2c
ffffd698
ffffd828
ffffd9b8
ffffdb48
ffffdcd8
ffffde68
ffffdff8
ffffe188
Segmentation fault
cnitfiftythree@deb:~/127/ch2$
```

- Cannot write to address ffffe188

Debug

Insert
breakpoint
and run

```
cnitfiftythree@deb:~/127/ch2$ gdb -q ch2c
Reading symbols from ch2c...done.
(gdb) list
warning: Source file is more recent than executable.
1      #include <stdio.h>
2
3      int main()
4      {
5          int array[5], i;
6          for (i=0; i<=1000; i+=100)
7          {
8              printf("%x\n", &array[i]);
9              array[i] = 1;
10         }
(gdb) break 9
Breakpoint 1 at 0x5e4: file ch2c.c, line 9.
(gdb) run
Starting program: /home/cnitfiftythree/127/ch2/ch2c
ffffd638

Breakpoint 1, main () at ch2c.c:9
9              array[i] = 1;
(gdb) █
```

Memory Map

```
(gdb) info proc mappings
process 6813
Mapped address spaces:

   Start Addr   End Addr       Size     Offset objfile
   0x56555000   0x56556000     0x1000      0x0  /home/cnitfiftythree/127/ch2/ch2c
   0x56556000   0x56557000     0x1000      0x0  /home/cnitfiftythree/127/ch2/ch2c
   0x56557000   0x56558000     0x1000     0x1000  /home/cnitfiftythree/127/ch2/ch2c
   0x56558000   0x56579000    0x21000      0x0  [heap]
   0xf7e14000   0xf7fc5000    0x1b1000      0x0  /lib32/libc-2.24.so
   0xf7fc5000   0xf7fc6000     0x1000     0x1b1000  /lib32/libc-2.24.so
   0xf7fc6000   0xf7fc8000     0x2000     0x1b1000  /lib32/libc-2.24.so
   0xf7fc8000   0xf7fc9000     0x1000     0x1b3000  /lib32/libc-2.24.so
   0xf7fc9000   0xf7fcc000     0x3000      0x0
   0xf7fd3000   0xf7fd5000     0x2000      0x0
   0xf7fd5000   0xf7fd7000     0x2000      0x0  [vvar]
   0xf7fd7000   0xf7fd9000     0x2000      0x0  [vdso]
   0xf7fd9000   0xf7ffc000    0x23000      0x0  /lib32/ld-2.24.so
   0xf7ffc000   0xf7ffd000     0x1000     0x22000  /lib32/ld-2.24.so
   0xf7ffd000   0xf7ffe000     0x1000     0x23000  /lib32/ld-2.24.so
   0xffffd000   0xffffe000    0x21000      0x0  [stack]
```

(gdb)

- Stack ends at 0xffffe000
- Trying to write past this address caused a segmentation fault

The Stack

LIFO (Last-In, First-Out)

- ESP (Extended Stack Pointer) register points to the top of the stack
- PUSH puts items on the stack
 - push 1
 - push addr var

Address Value
643410h Address of variable VAR
643414h 1
643418h

← ESP points to this address

Stack

- POP takes items off the stack
 - pop eax
 - pop ebx

Address	Value
643410h	Address of variable VAR
643414h	1
643418h	

← ESP points to this address

EBP (Extended Base Pointer)

- EBP is typically used for calculated addresses on the stack
 - `mov eax, [ebp+10h]`
- Copies the data 16 bytes down the stack into the EAX register

Functions and the Stack

Purpose

- The stack's primary purpose is to make the use of functions more efficient
- When a function is called, these things occur:
 - Calling routine stops processing its instructions
 - Saves its current state
 - Transfers control to the function
 - Function processes its instructions
 - Function exits
 - State of the calling function is restored
 - Calling routine's execution resumes

Figure 2-3: Visual representation of the stack after a function has been called

Array
EBP
RET
A
B

Low Memory Addresses and Top of the Stack

High Memory Addresses and Bottom of the Stack

Functions and the Stack


- Primary purpose of the stack
 - To make functions more efficient
- When a function is called
 - Push function's **arguments** onto the stack
 - Call function, which pushes the return address **RET** onto the stack, which is the **EIP** at the time the function is called

Functions and the Stack

- Before function starts, a **prolog** executes, pushing **EBP** onto the stack
- It then copies **ESP** into **EBP**
- Calculates size of local variables
- Reserves that space on the stack, by subtracting the size from **ESP**
- Pushes local variables onto stack

Functions and the Stack

Low memory
addresses &
top of Stack



High memory
addresses &
bottom of Stack

```
#include <stdio.h>

void function(int a, int b)
{
    int array[5];
}

main()
{
    function(1,2);
    printf("This is where the  
return address points\n");
}
```


Example of a Function

```
GNU nano 2.9.2 ch2d.c

#include <stdio.h>

void function(int a, int b)
{
    int array[5] = {1, 2, 3, 4, 5};
}

int main()
{
    function(1,2);
    printf("Returned from function\n");
}
```

```
cnitfiftythree@deb:~/127/ch2$ gcc -m32 -g -o ch2d ch2d.c
cnitfiftythree@deb:~/127/ch2$ ./ch2d
Returned from function
cnitfiftythree@deb:~/127/ch2$
```

Debug and Set Breakpoints

```
cnitfiftythree@deb:~/127/ch2$ gdb -q ch2d
Reading symbols from ch2d...done.
(gdb) list
1      #include <stdio.h>
2
3      void function(int a, int b)
4      {
5          int array[5] = {1, 2, 3, 4, 5};
6      }
7
8      int main()
9      {
10         function(1, 2);
(gdb) break 10
Breakpoint 1 at 0x5e0: file ch2d.c, line 10.
(gdb) break 6
Breakpoint 2 at 0x5c3: file ch2d.c, line 6.
(gdb) █
```

In main()

```
(gdb) run
Starting program: /home/cnitfiftythree/127/ch2/ch2d

Breakpoint 1, main () at ch2d.c:10
10          function(1, 2);
(gdb) info registers
eax                0xf7fc9dbc          -134439492
ecx                0xfffffd670        -10640
edx                0xfffffd694        -10604
ebx                0x56557000          1448439808
esp                0xfffffd650        0xfffffd650
ebp                0xfffffd658        0xfffffd658
esi                0x1                1
edi                0xf7fc8000          -134447104
eip                0x565555e0          0x565555e0 <main+26>
```

Stack frame goes from ebp to esp

In function()

```
(gdb) continue
Continuing.

Breakpoint 2, function (a=1, b=2) at ch2d.c:6
6      }
(gdb) info registers
eax                0x56557000          1448439808
ecx                0xfffffd670        -10640
edx                0xfffffd694        -10604
ebx                0x56557000          1448439808
esp                0xfffffd620        0xfffffd620
ebp                0xfffffd640        0xfffffd640
esi                0x1                1
edi                0xf7fc8000         -134447104
eip                0x565555c3         0x565555c3 <function+51>
```

Stack frame goes from ebp to esp

Examine the Stack Frame

```
(gdb) info registers
eax          0x56557000      1448439808
ecx          0xffffd670      -10640
edx          0xffffd694      -10604
ebx          0x56557000      1448439808
esp          0xffffd620      0xffffd620
ebp          0xffffd640      0xffffd640
esi          0x1             1
edi          0xf7fc8000      -134447104
eip          0x565555c3      0x565555c3 <function+51>
eflags       0x216          [ PF AF IF ]
cs           0x23           35
ss           0x2b           43
ds           0x2b           43
es           0x2b           43
fs           0x0            0
gs           0x63           99
(gdb) x/20x $esp
0xffffd620: 0xf7fc8000 0xffffd704 0xf7ffcd00 0x00000001
0xffffd630: 0x00000002 0x00000003 0x00000004 0x00000005
0xffffd640: 0xffffd658 0x565555e9 0x00000001 0x00000002
0xffffd650: 0xffffd670 0x00000000 0x00000000 0xf7e2c286
0xffffd660: 0x00000001 0xf7fc8000 0x00000000 0xf7e2c286
(gdb)
```

- Highlighted region is the **stack frame of function()**
- Below it is the stack frame of **main()**

Disassemble Main

Dump of assembler code for function main:

```
0x565555c6 <+0>:    lea     0x4(%esp),%ecx
0x565555ca <+4>:    and     $0xfffffffff0,%esp
0x565555cd <+7>:    pushl   -0x4(%ecx)
0x565555d0 <+10>:   push    %ebp
0x565555d1 <+11>:   mov     %esp,%ebp
0x565555d3 <+13>:   push    %ebx
0x565555d4 <+14>:   push    %ecx
0x565555d5 <+15>:   call    0x56555460 <__x86.get_pc_thunk.bx>
0x565555da <+20>:   add     $0x1a26,%ebx
0x565555e0 <+26>:   push    $0x2
0x565555e2 <+28>:   push    $0x1
0x565555e4 <+30>:   call    0x56555590 <function>
0x565555e9 <+35>:   add     $0x8,%esp
```

- To call a function:
 - push arguments onto the stack
 - call the function

Disassemble Function

```
(gdb) disassemble function
Dump of assembler code for function function:
0x56555590 <+0>:      push    %ebp
0x56555591 <+1>:      mov     %esp,%ebp
0x56555593 <+3>:      sub     $0x20,%esp
0x56555596 <+6>:      call    0x5655560d <__x86.get_pc_thunk.ax>
0x5655559b <+11>:     add     $0x1a65,%eax
0x565555a0 <+16>:     movl    $0x1,-0x14(%ebp)
0x565555a7 <+23>:     movl    $0x2,-0x10(%ebp)
0x565555ae <+30>:     movl    $0x3,-0xc(%ebp)
0x565555b5 <+37>:     movl    $0x4,-0x8(%ebp)
0x565555bc <+44>:     movl    $0x5,-0x4(%ebp)
=> 0x565555c3 <+51>:     nop
0x565555c4 <+52>:     leave
0x565555c5 <+53>:     ret
End of assembler dump.
(gdb) █
```

- Prolog:
 - **push** ebp onto stack
 - **mov** esp into ebp, starting a new stack frame
 - **sub** from esp, reserving room for local variables

Saved Return Address

```
(gdb) x/20x $esp
0xffffd620: 0xf7fc8000  0xffffd704  0xf7ffcd00  0x00000001
0xffffd630: 0x00000002  0x00000003  0x00000004  0x00000005
0xffffd640: 0xffffd658  0x565555e9  0x00000001  0x00000002
0xffffd650: 0xffffd670  0x00000000  0x00000000  0xf7e2c286
0xffffd660: 0x00000001  0xf7fc8000  0x00000000  0xf7e2c286
```

- Next word after stack frame
- Address of next instruction to be executed in main()

```
Dump of assembler code for function main:
0x565555c6 <+0>:  lea    0x4(%esp),%ecx
0x565555ca <+4>:  and    $0xffffffff0,%esp
0x565555cd <+7>:  pushl  -0x4(%ecx)
0x565555d0 <+10>: push    %ebp
0x565555d1 <+11>: mov     %esp,%ebp
0x565555d3 <+13>: push    %ebx
0x565555d4 <+14>: push    %ecx
0x565555d5 <+15>: call    0x56555460 <__x86.get_pc
0x565555da <+20>: add     $0x1a26,%ebx
0x565555e0 <+26>: push    $0x2
0x565555e2 <+28>: push    $0x1
0x565555e4 <+30>: call    0x56555590 <function>
0x565555e9 <+35>: add     $0x8,%esp
```


Stack Buffer Overflow Exploit

Stack Buffer Overflow Vulnerability

```
GNU nano 2.9.2 ch2e.c

#include <stdio.h>

void user_input(void)
{
    char buf[30];
    gets(buf);
    printf("%s\n", buf);
}

int main()
{
    user_input();
    return 0;
}
```

gets() reads user input

Does not limit its length

Compile and Run

```
cnitfiftythree@deb:~/127/ch2$ gcc -m32 -g -o ch2e ch2e.c
ch2e.c: In function 'user_input':
ch2e.c:6:2: warning: implicit declaration of function 'gets' [-Wimplicit-function-declaration]
     gets(buf);
     ^~~~
/tmp/ccMhxXHA.o: In function 'user_input':
/home/cnitfiftythree/127/ch2/ch2e.c:6: warning: the 'gets' function is dangerous and should not be used.
cnitfiftythree@deb:~/127/ch2$ ./ch2e
HELLO
HELLO
cnitfiftythree@deb:~/127/ch2$ ./ch2e
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
Segmentation fault
cnitfiftythree@deb:~/127/ch2$
```

Segmentation fault indicates an illegal operation

Debug and Set Breakpoint

```
cnitfiftythree@deb:~/127/ch2$ gdb -q ch2e
Reading symbols from ch2e...done.
(gdb) list
1      #include <stdio.h>
2
3      void user_input(void)
4      {
5          char buf[30];
6          gets(buf);
7          printf("%s\n", buf);
8      }
9
10     int main()
(gdb) break 7
Breakpoint 1 at 0x5e1: file ch2e.c, line 7.
(gdb) █
```

Break after **gets()**

Stack After HELLO

```
(gdb) run
Starting program: /home/critfiftythree/127/ch2/ch2e
HELLO

Breakpoint 1, user_input () at ch2e.c:7
7       printf("%s\n", buf);
(gdb) x/20x $esp
0xffffd620: 0x4548800c 0x004f4c4e 0xf7ffcd00 0x00040000
0xffffd630: 0x00000000 0x56557000 0x00000001 0x5655567b
0xffffd640: 0x00000001 0x00000000 0xffffd658 0x56555616
0xffffd650: 0xf7fc83dc 0xffffd670 0x00000000 0xf7e2c286
0xffffd660: 0x00000001 0xf7fc8000 0x00000000 0xf7e2c286
(gdb) info registers
eax          0xffffd622      -10718
ecx          0xfbad2288     -72539512
edx          0xf7fc987c     -134440836
ebx          0x56557000     1448439808
esp          0xffffd620     0xffffd620
ebp          0xffffd648     0xffffd648
esi          0x1            1
edi          0xf7fc8000     -134447104
eip          0x565555e1     0x565555e1 <user_input+33>
```

- ASCII values for **HELLO** appear in the words outlined in red
- Return value is outlined in green

ASCII

- Google "ASCII"
 - 0x41 is A
 - 0x42 is B
 - etc.

Dec	Hx	Oct	Html	Chr
64	40	100	@	@
65	41	101	A	A
66	42	102	B	B
67	43	103	C	C
68	44	104	D	D
69	45	105	E	E
70	46	106	F	F
71	47	107	G	G
72	48	110	H	H
73	49	111	I	I
74	4A	112	J	J
75	4B	113	K	K
76	4C	114	L	L
77	4D	115	M	M
78	4E	116	N	N
79	4F	117	O	O
80	50	120	P	P
81	51	121	Q	Q
82	52	122	R	R

Stack After AAAAAA...

```
(gdb) run
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /home/cniftythree/127/ch2/ch2e
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
Breakpoint 1, user_input () at ch2e.c:7
7         printf("%s\n", buf);
(gdb) x/20x $esp
0xffffd620: 0x41418000  0x41414141  0x41414141  0x41414141
0xffffd630: 0x41414141  0x41414141  0x41414141  0x41414141
0xffffd640: 0x41414141  0x41414141  0x41414141  0x41414141
0xffffd650: 0x41414141  0x41414141  0x41414141  0xf7004141
0xffffd660: 0x00000001  0xf7fc8000  0x00000000  0xf7e2c286
(gdb)
```

- Stack frame is filled with many A characters
- Return value is overwritten with 0x41414141

Examining the Crash

```
(gdb) continue
Continuing.
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
Program received signal SIGSEGV, Segmentation fault.
0x41414141 in ?? ()
(gdb) info registers
eax                0x3d          61
ecx                0xfbad0084      -72548220
edx                0xf7fc9870      -134440848
ebx                0x41414141      1094795585
esp                0xfffffd650     0xfffffd650
ebp                0x41414141      0x41414141
esi                0x1           1
edi                0xf7fc8000      -134447104
eip                0x41414141      0x41414141
eflags             0x10282    [ SF IF RF ]
```

- eip value is 0x41414141
- Controlled by user input!

gdb Commands

list	<i>show source code</i>
run	<i>execute program</i>
break	<i>insert breakpoint</i>
x	<i>examine memory</i>
disassemble	<i>show assembly code</i>
continue	<i>resume execution</i>
info registers	<i>see registers</i>
info proc mapping	<i>see memory map</i>

Kahoot!